



Big Data

Tutorial #4

Isabelle Kuhlmann

2020-05-29



Outline

- 1) Lecture Recap
- 2) Exercise Discussion
- 3) Joins in Spark
- 4) Communication Costs
- 5) Homework Exercise

What previously happened...

Lecture Recap

Broadcast Join & Shuffle Join

- Broadcast join:
 - If one table is small (i.e., if it fits completely into the memory of each executor), collect it at the driver
 - Broadcast the complete table to all executors
 - Communication between driver and executors only at the beginning
- Shuffle join:
 - If both tables are big
 - Shuffle: communication between executors
 - Every node talks to every other node
 - Data sharing with respect to keys

Cost Measures for Algorithms

- Two main costs:
 - **Communication costs**, i.e., what does it cost to send data around the network?
 - Measured, for instance, in bytes
 - **Computation costs**, i.e., what are the costs for mappers, reducers, and the system?
 - Measured, for instance, in hours of computing power
- *Communication cost*:
 - Total I/O of all processes
- *Elapsed communication cost*:
 - Max of I/O along any path
- *(Elapsed) computation cost*:
 - Analogous, but count only running time of processes

Efficiency Aspects of Spark

- Volcano Iterator Model
 - Queries are transformed into sequences of operations
- Whole-stage code generation
 - Works well with simple operations
- Vectorization
 - Apply an instruction to an entire array
 - Must be supported by CPU

Retail data

Exercise Discussion

Homework Assignment – Discussion

- Consider the retail dataset again. Answer the following questions in two different ways: using SQL and using Spark code (DataFrame API).
 - How many orders did customers perform at which hour?
 - How frequently was each product bought in the different countries?
- Demo

Natural join, inner join, semi-join, etc.

Joins in Spark

Different Types of Joins

- Natural join
- Inner join
- (Left/right) outer join
- Left semi-join
- Left anti-join
- Cross join (cartesian product)

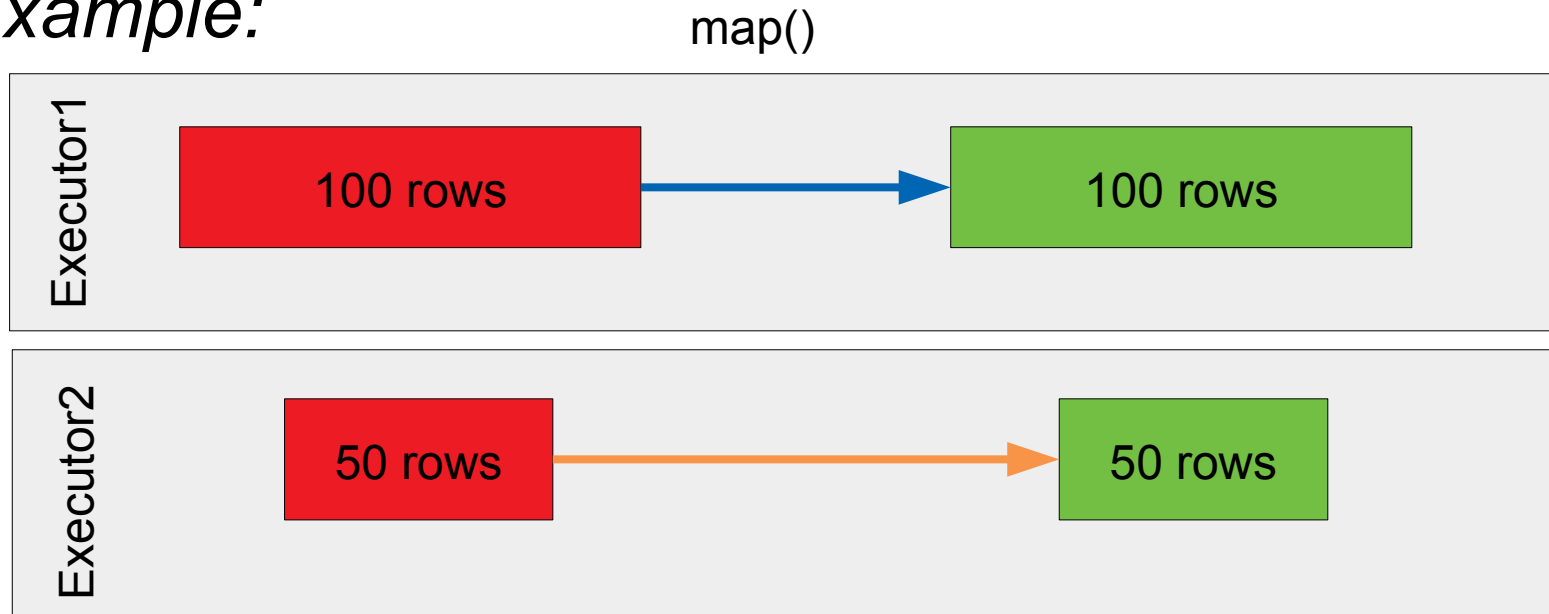
➔ Demo

Considering narrow and wide transformations

Communication Costs

Costs for Narrow Transformation

- Example:*




- Communication cost?
- Elapsed communication cost?

Costs for Narrow Transformation

- *Example:*

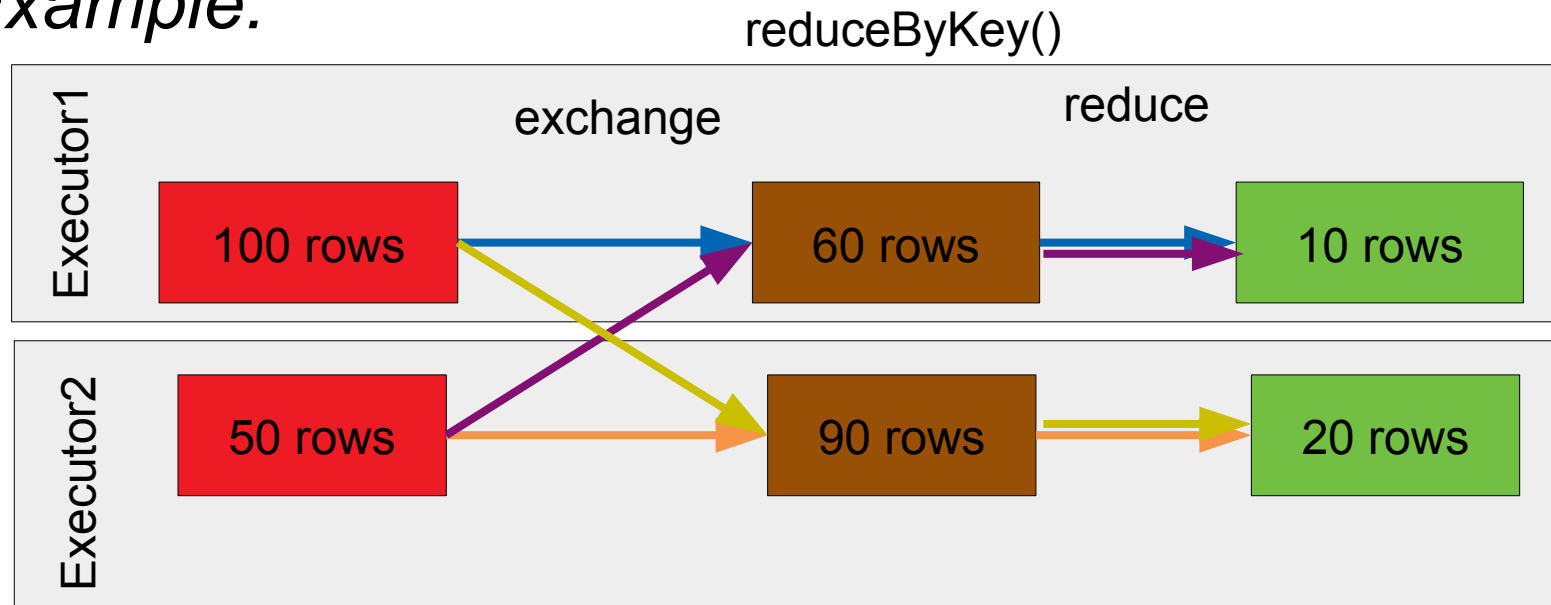
map()



- Communication cost
 $= 100 + 100 + 50 + 50 = 300$
- Elapsed communication cost
 $= \max((100 + 100), (50 + 50)) = 200$

Costs for Wide Transformations (Shuffle)

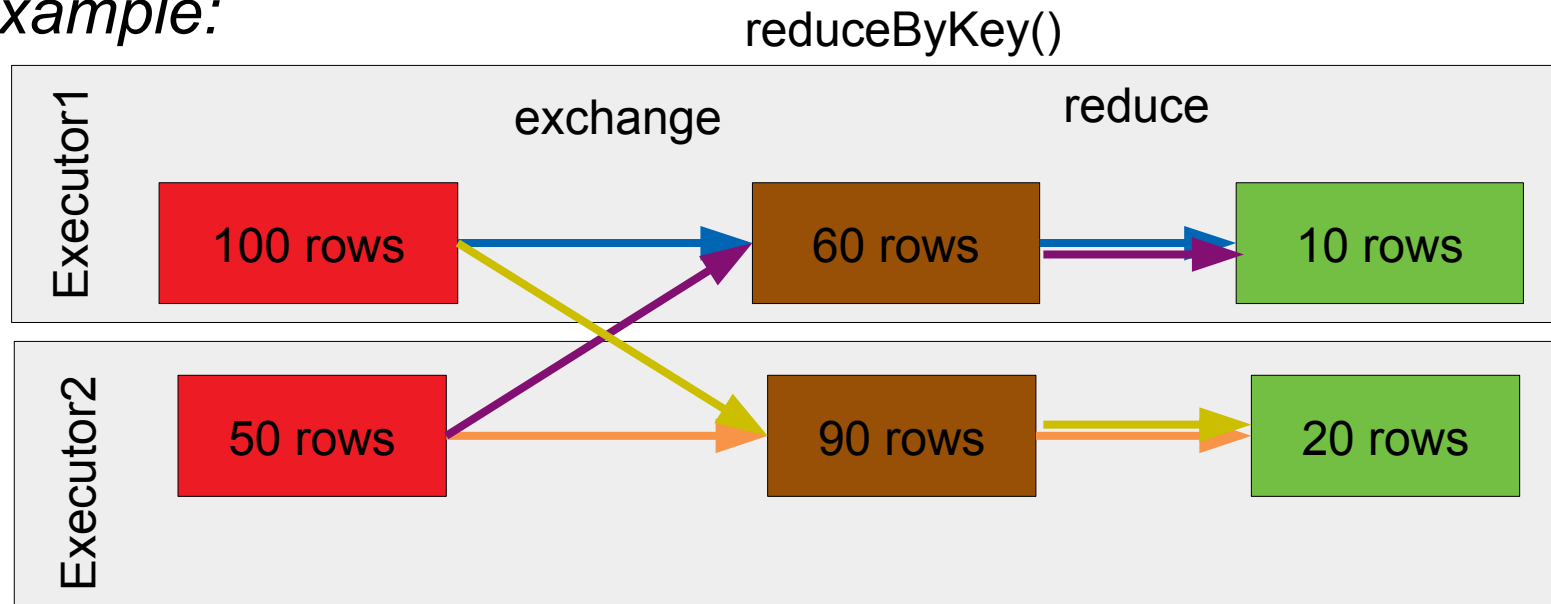
- *Example:*



- Communication cost?
- Elapsed communication cost?

Costs for Wide Transformations (Shuffle)

- *Example:*



- **Communication cost**
 = exchange cost + reduce cost
 = $[100 + 50 + 60 + 90] + [60 + 90 + 10 + 20]$
 = 480
- **Elapsed communication cost**
 = $\max(100+2*60+10, 100+2*90+20, 50+2*60+10, 50+2*90+20)$
 = $\max(230, 300, 180, 250)$
 = 300

Now it's your turn!

Homework Exercise

Exercise #1 – Spark

1) Pivot

Consider the *Retail dataset* again.

How many instances of each product were sold in each country?

2) Joins

Just play around with the different join types and make sure you understand each one.

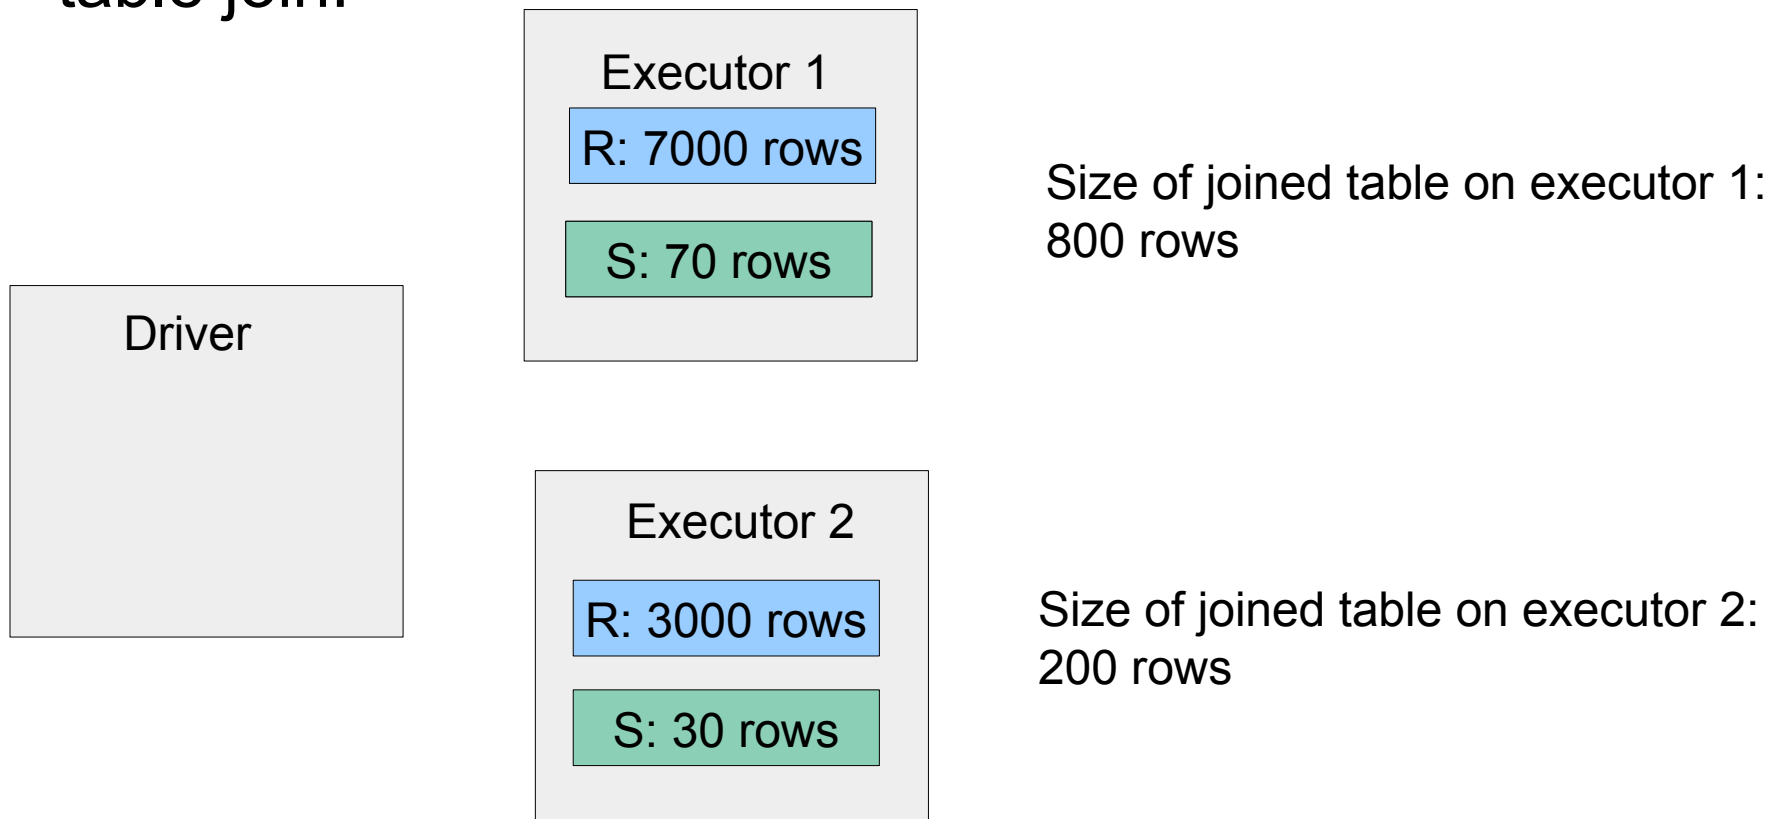
You can utilize `customers.csv` and `orders(2).csv` for this purpose.

➔ These files can be found in OLAT under:

`Tutorial Material/Tutorial04`

Exercise #2 – Communication Cost

- Compute the **communication cost** and the **elapsed communication cost** for the following big table-to-small table join:





Thank you for your Attention!

